

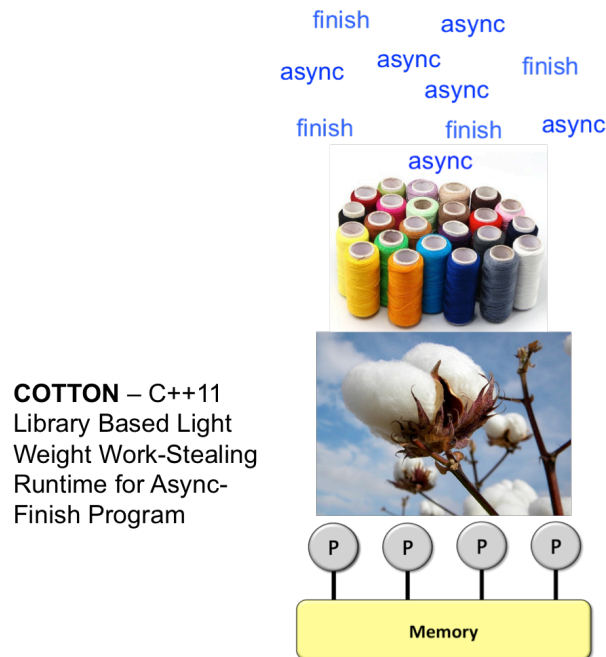
Homework Assignment-2

Due by 11:59pm on 04th February 2018

(Total 5% weightage)

Instructor: Vivek Kumar

COTTON: A light-weight work-stealing runtime for async-finish parallelism



No extensions will be provided. Any submission after the deadline will not be evaluated. If you see an ambiguity or inconsistency in a question, please seek a clarification from the teaching staff.

Plagiarism: This is a pair programming based assignment that you must do with the group member that you have already chosen. No change to group is allowed. You are not allowed to discuss the approach/solution outside your group. You should never misrepresent some other group's work as your own. In case any plagiarism case is detected, it will be dealt as per the new plagiarism policy of IITD and will be applied to each member in the group.

Instructions

- 1) Hardware requirements:
 - a. You will require a Linux/Mac OS to complete this assignment.
 - b. You can easily code/complete this assignment on your laptop.
- 2) Software requirements are:


```

cotton::finalize_runtime();           // finalize and release runtime resources
return 0;
}

```

Cotton should support very basic async-finish parallelism. It should be able to support “flat finish” scopes. A flat-finish scope is a single-level finish scope for all async tasks spawned within its scope (i.e., an async task can recursively create new async tasks but it can never spawn new finish scopes).

Runtime Implementation

Cotton runtime should be a library-based **thread pool** based implementation that would rely on “**work-stealing**” technique for task scheduling. The rough code sketch for this type of runtime implementation was explained in Lecture 07 and Lecture 08. You are encouraged to use that code but remember that was pseudo-code and for actual implementation you will have to fill up the missing parts.

Requirements:

- a) You should implement cotton in C++.
- b) Proper documentation should be done in your coding.
- c) All runtime APIs should be within the “cotton” namespace (see pseudo-code above).
- d) As explained in Lecture 08, your work-stealing runtime should have per-thread deque datastructure. To access this deque you will need task counters at both the ends of the deque. The bottom (or tail) will be private to victim (owner) and will be used in LIFO order. The top (or head) will only be accessed by a thief during steal operation (FIFO order).
- e) To implement concurrent deque accesses you should **only** use `pthread_mutex_lock/pthread_mutex_unlock` APIs. You won’t be able to the lockless deque explained in lectures without using memory barriers and atomic compare-and-swap operations that are not covered in lectures. Hence, in your locked deque implementation the `push()` operation will always be lockless whereas `steal()` and `pop()` would always be synchronized.
- f) The size of the deque should be a compile time constant.
- g) Your cotton runtime should throw proper error if deque becomes full. In such scenarios, increasing the deque size (a compile time constant as explained above), rebuilding cotton runtime and your testcase should resolve this issue.
- h) Cotton users should be able to control cotton’s thread pool size by using the environment variable “COTTON_WORKERS” (e.g. usage: `export COTTON_WORKERS=4`). Default thread pool size should be 1. The master thread should also be a part of the thread pool.

Your assignment directory should compose of following files **ONLY** (don't change the file names and its functionality):

- a. **cotton.h**: This header file should only contain the declarations of `async`, `start_finish`, `end_finish`, `init_runtime`, and `finalize_runtime` (see the pseudo-code for input program above).
- b. **cotton-runtime.cpp**: This should contain all of your cotton runtime implementation.
- c. **cotton-runtime.h**: This header file should contain the declarations of all the routines you implemented inside `cotton-runtime.cpp`.
- d. **Makefile**: This should be able to compile your implementation of cotton as well as the standard testcase that we are providing herewith (see below).

We will evaluate your implementation of cotton using the NQueens application (with input **N=12**). A parallel version of NQueens application is provided to you that is written using cotton APIs. You should not do any changes to this application as for evaluation the instructor will use his own copy of NQueens. This is the same copy that you have been provided along with this assignment. We will consider the correctness of your implementation only if NQueens executes successfully for any value of thread pool size between 1 and 4 (inclusive). Your cotton runtime should cleanly exit after NQueens completes. An example of NQueens execution using cotton is as following:

```
COTTON_WORKERS=4 ./NQueens 12
```

Your assignment submission should consist of two parts:

- a) A tar file containing your source files as mentioned above in the requirements.
- b) Your writeup (as PDF inside the above tar) that explains the following:
 - a. The speedup graph of execution NQueens with threads=1,2,3, and 4; with input as 12.
 - b. **Listing individual contribution from each group members (you won't get any marks if you miss this part).**