

Paper 5: Energy-Efficient Work-Stealing Language Runtimes

Review

Dattatreya Mohapatra¹ Viraj Parimi²

¹2015021

²2015068

FPP, Winter 2018

1 Background

- DVFS
- Work-Stealing

2 Problem Statement

3 Related Work

4 Contributions

5 Motivation

6 Insights

- Workpath-Sensitive TC
- Workload-Sensitive TC
- Unified Algorithm

7 Implementation

- Order of Immediacy
- Tempo-Frequency Mapping
- Worker-Core Mapping

- Tempo Setting at Idle
- Overheads

8 Experimental Methodology

- Benchmarks
- System Specs
- Energy Consumption

9 Results

- Overall Results
- Relative Effectiveness
- Effect of Frequency Selection
- N-Frequency Tempo Control
- Static vs Dynamic Scheduling
- Naive Frequency Scaling

10 Conclusion

11 COTTON Integration

Outline

1 Background

- DVFS
- Work-Stealing

2 Problem Statement

3 Related Work

4 Contributions

5 Motivation

6 Insights

7 Implementation

8 Experimental Methodology

9 Results

10 Conclusion

11 COTTON Integration

Outline

1 Background

- DVFS
- Work-Stealing

2 Problem Statement

3 Related Work

4 Contributions

5 Motivation

6 Insights

7 Implementation

8 Experimental Methodology

9 Results

10 Conclusion

11 COTTON Integration

- Dynamic Voltage and Frequency Scaling (DVFS) – a framework to change the frequency and/or operating voltage of a processor(s) based on system requirements.

- Dynamic Voltage and Frequency Scaling (DVFS) – a framework to change the frequency and/or operating voltage of a processor(s) based on system requirements.
- Uses CPUFreq, a linux kernel framework, that decides on whether to increase/decrease frequency of a processor(s).

- Dynamic Voltage and Frequency Scaling (DVFS) – a framework to change the frequency and/or operating voltage of a processor(s) based on system requirements.
- Uses CPUFreq, a linux kernel framework, that decides on whether to increase/decrease frequency of a processor(s).
- Consists of two elements:
 - Governor : Directs the driver to change frequency when required.
 - Driver : Performs actions based on governor's decision.

Outline

1 Background

- DVFS
- Work-Stealing

2 Problem Statement

3 Related Work

4 Contributions

5 Motivation

6 Insights

7 Implementation

8 Experimental Methodology

9 Results

10 Conclusion

11 COTTON Integration

Background

Work-Stealing

- Work-Stealing
 - Work-First Principle
 - Deque Management
 - Work-Stealing Scheduler
- Covered in course lectures.

Outline

- 1 Background
- 2 Problem Statement**
- 3 Related Work
- 4 Contributions
- 5 Motivation
- 6 Insights
- 7 Implementation
- 8 Experimental Methodology
- 9 Results
- 10 Conclusion
- 11 COTTON Integration

Problem Statement

Objective

Improve energy efficiency of work stealing runtime with minimal performance loss

Outline

- 1 Background
- 2 Problem Statement
- 3 Related Work
- 4 Contributions
- 5 Motivation
- 6 Insights
- 7 Implementation
- 8 Experimental Methodology
- 9 Results
- 10 Conclusion
- 11 COTTON Integration

Related Work

- Most of the related work can be summarized in two more established areas.

Related Work

- Most of the related work can be summarized in two more established areas.
- Optimization of work-stealing runtimes
 - *A-Steal*: An adaptive scheduler to take parallelism feedback into account.
 - *SLAW*: Adds adaptive scheduling policies based on locality information.
 - *AdaptiveTC*: Improves system performance through adaptive thread management.
 - *BWS*: Improves system throughput and fairness in time-sharing multicore systems.

Related Work

- Most of the related work can be summarized in two more established areas.
- Optimization of work-stealing runtimes
 - *A-Steal*: An adaptive scheduler to take parallelism feedback into account.
 - *SLAW*: Adds adaptive scheduling policies based on locality information.
 - *AdaptiveTC*: Improves system performance through adaptive thread management.
 - *BWS*: Improves system throughput and fairness in time-sharing multicore systems.
- Energy efficiency of multi-threaded programs
 - Various DVFS-based solutions have been proposed earlier.
 - *Magklis et. al.* designed a profiling-based DVFS algorithm on CPUs with multiple clock domains.
 - *Wu et. al.* designed a DVFS-based strategy where the interval of DVFS use is adaptive to recent instance issue queue occupancy.

Related Work

- Most of the related work can be summarized in two more established areas.
- Optimization of work-stealing runtimes
 - *A-Steal*: An adaptive scheduler to take parallelism feedback into account.
 - *SLAW*: Adds adaptive scheduling policies based on locality information.
 - *AdaptiveTC*: Improves system performance through adaptive thread management.
 - *BWS*: Improves system throughput and fairness in time-sharing multicore systems.
- Energy efficiency of multi-threaded programs
 - Various DVFS-based solutions have been proposed earlier.
 - *Magklis et. al.* designed a profiling-based DVFS algorithm on CPUs with multiple clock domains.
 - *Wu et. al.* designed a DVFS-based strategy where the interval of DVFS use is adaptive to recent instance issue queue occupancy.
- Other approaches include thread migration, hardware/software approximation, or a combination of all of them.

Outline

- 1 Background
- 2 Problem Statement
- 3 Related Work
- 4 Contributions**
- 5 Motivation
- 6 Insights
- 7 Implementation
- 8 Experimental Methodology
- 9 Results
- 10 Conclusion
- 11 COTTON Integration

- HERMES – The first framework addressing energy efficiency in work-stealing systems.

- HERMES – The first framework addressing energy efficiency in work-stealing systems.
- Programs are *tempo-enabled* : Different threads may execute at different speeds(*tempo*). Two novel, complementary tempo control strategies:
 - workpath-sensitive
 - workload-sensitive

- HERMES – The first framework addressing energy efficiency in work-stealing systems.
- Programs are *tempo-enabled* : Different threads may execute at different speeds(*tempo*). Two novel, complementary tempo control strategies:
 - workpath-sensitive
 - workload-sensitive
- A prototyped implementation demonstrating an average of 11-12% energy savings with 3-4% performance loss over work-stealing benchmarks.

Outline

- 1 Background
- 2 Problem Statement
- 3 Related Work
- 4 Contributions
- 5 Motivation**
- 6 Insights
- 7 Implementation
- 8 Experimental Methodology
- 9 Results
- 10 Conclusion
- 11 COTTON Integration

Motivation

- In the multi-core era, work stealing received considerable interest in language runtime design.

- In the multi-core era, work stealing received considerable interest in language runtime design.
- There is an active interest in research improving its performance-critical properties, such as adaptiveness, scalability, and fairness.

- In the multi-core era, work stealing received considerable interest in language runtime design.
- There is an active interest in research improving its performance-critical properties, such as adaptiveness, scalability, and fairness.
- In comparison, energy efficiency in work-stealing systems has received little attention, which is particularly unfortunate.

Outline

1 Background

2 Problem Statement

3 Related Work

4 Contributions

5 Motivation

6 Insights

- Workpath-Sensitive TC
- Workload-Sensitive TC
- Unified Algorithm

7 Implementation

8 Experimental Methodology

9 Results

10 Conclusion

11 COTTON Integration

Outline

1 Background

2 Problem Statement

3 Related Work

4 Contributions

5 Motivation

6 Insights

- Workpath-Sensitive TC
- Workload-Sensitive TC
- Unified Algorithm

7 Implementation

8 Experimental Methodology

9 Results

10 Conclusion

11 COTTON Integration

Insights

Workpath-Sensitive Tempo Control

- Determines thread tempo based on control flow, with threads tackling *immediate work* executing at a faster tempo.

- Determines thread tempo based on control flow, with threads tackling *immediate work* executing at a faster tempo.
- Victim worker takes precedence over the thief worker in a thief-victim relationship.

- Determines thread tempo based on control flow, with threads tackling *immediate work* executing at a faster tempo.
- Victim worker takes precedence over the thief worker in a thief-victim relationship.
- Two important design ideas:
 - Thief Procrastination
 - Immediacy Relay

Thief Procrastination

At the beginning of the thief-victim relationship, the tempo of the thief worker should be set to be slower than the victim worker.

Immediacy Relay

If the thief-victim relationship terminates because the victim runs out of work, the tempo of all the thieves should be raised. This happens in a recursive fashion.

Outline

1 Background

2 Problem Statement

3 Related Work

4 Contributions

5 Motivation

6 Insights

- Workpath-Sensitive TC
- **Workload-Sensitive TC**
- Unified Algorithm

7 Implementation

8 Experimental Methodology

9 Results

10 Conclusion

11 COTTON Integration

- Selects the appropriate thread tempo based on the size of work-stealing deques.

Insights

Workload-Sensitive Tempo Control

- Selects the appropriate thread tempo based on the size of work-stealing dequeues.
- Threads with longer dequeues run at faster tempo.

Insights

Workload-Sensitive Tempo Control

- Selects the appropriate thread tempo based on the size of work-stealing deque.
- Threads with longer deques run at faster tempo.
- Threshold calculation:

$$thld_i = \left(\frac{2 \times L}{K + 1} \right) \times i$$

- L = average deque size
- K = number of thresholds
- $i \in [1, K]$

Outline

1 Background

2 Problem Statement

3 Related Work

4 Contributions

5 Motivation

6 Insights

- Workpath-Sensitive TC
- Workload-Sensitive TC
- **Unified Algorithm**

7 Implementation

8 Experimental Methodology

9 Results

10 Conclusion

11 COTTON Integration

- In workpath-alone executions, a non-immediate worker may have many deque items to work on.

- In workpath-alone executions, a non-immediate worker may have many deque items to work on.
- Similarly, in workload-alone executions, an immediate worker with fewer deque items is set to a lower tempo.

- In workpath-alone executions, a non-immediate worker may have many deque items to work on.
- Similarly, in workload-alone executions, an immediate worker with fewer deque items is set to a lower tempo.
- Increased workpath length means increased execution time, which potentially leads to additional energy consumption.

- In workpath-alone executions, a non-immediate worker may have many deque items to work on.
- Similarly, in workload-alone executions, an immediate worker with fewer deque items is set to a lower tempo.
- Increased workpath length means increased execution time, which potentially leads to additional energy consumption.
- In both cases, a second opinion from the other strategy can lead to better decisions.

- 1 Background
- 2 Problem Statement
- 3 Related Work
- 4 Contributions
- 5 Motivation
- 6 Insights
- 7 Implementation
 - Order of Immediacy
 - Tempo-Frequency Mapping
 - Worker-Core Mapping
 - Tempo Setting at Idle
 - Overheads
- 8 Experimental Methodology
- 9 Results
- 10 Conclusion
- 11 COTTON Integration

- 1 Background
- 2 Problem Statement
- 3 Related Work
- 4 Contributions
- 5 Motivation
- 6 Insights
- 7 **Implementation**
 - Order of Immediacy
 - Tempo-Frequency Mapping
 - Worker-Core Mapping
 - Tempo Setting at Idle
 - Overheads
- 8 Experimental Methodology
- 9 Results
- 10 Conclusion
- 11 COTTON Integration

Implementation

Order of Immediacy

- Key data-structure for workpath sensitivity.
- Double linked list across workers connected by `next` and `prev` pointers.
- When a victim terminates, immediacy relay happens through this linked list.

Implementation

Order of Immediacy

```
structure WORKER  
  DQ // deque (array)  
  H // head index  
  T // tail index  
end structure
```

```
structure WORKER  
  DQ // deque (array)  
  H // head index  
  T // tail index  
  next // next immediate work  
  prev // prev immediate work  
  thld // size thresholds (array)  
  S // size threshold index  
end structure
```

- 1 Background
- 2 Problem Statement
- 3 Related Work
- 4 Contributions
- 5 Motivation
- 6 Insights
- 7 **Implementation**
 - Order of Immediacy
 - **Tempo-Frequency Mapping**
 - Worker-Core Mapping
 - Tempo Setting at Idle
 - Overheads
- 8 Experimental Methodology
- 9 Results
- 10 Conclusion
- 11 COTTON Integration

Implementation

Tempo-Frequency Mapping

- Achieved through DVFS as modern CPUs support a limited and discrete set of frequencies.

Implementation

Tempo-Frequency Mapping

- Achieved through DVFS as modern CPUs support a limited and discrete set of frequencies.
- Let $\{f_1, f_2, \dots, f_n\}$ be set of frequencies supported by the CPU core, where $f_i > f_{i+1}$ for any $i \in [1, n - 1]$

```
procedure DOWN(w, v)
     $f \leftarrow$  frequency of core hosting v
    if  $f == f_i$  and  $i < N$  then
        ...// scale core hosting w to  $f_{i+1}$ 
    end if
end procedure
```

Implementation

Tempo-Frequency Mapping

- Achieved through DVFS as modern CPUs support a limited and discrete set of frequencies.
- Let $\{f_1, f_2, \dots, f_n\}$ be set of frequencies supported by the CPU core, where $f_i > f_{i+1}$ for any $i \in [1, n - 1]$

```
procedure DOWN(w, v)
     $f \leftarrow$  frequency of core hosting v
    if  $f == f_i$  and  $i < N$  then
        ...// scale core hosting w to  $f_{i+1}$ 
    end if
end procedure
```

- $N \leq n$. This is called **N -frequency tempo control**.

1 Background

2 Problem Statement

3 Related Work

4 Contributions

5 Motivation

6 Insights

7 **Implementation**

- Order of Immediacy
- Tempo-Frequency Mapping
- **Worker-Core Mapping**
- Tempo Setting at Idle
- Overheads

8 Experimental Methodology

9 Results

10 Conclusion

11 COTTON Integration

Implementation

Worker-Core Mapping

- Relationship between workers (threads) and hosting CPU cores should be known.

Implementation

Worker-Core Mapping

- Relationship between workers (threads) and hosting CPU cores should be known.
- Two scheduling strategies are used.
 - Static : each worker is pre-assigned to CPU core.
 - Dynamic : each worker thread may migrate from one core to another during program execution.

Implementation

Worker-Core Mapping

- Relationship between workers (threads) and hosting CPU cores should be known.
- Two scheduling strategies are used.
 - Static : each worker is pre-assigned to CPU core.
 - Dynamic : each worker thread may migrate from one core to another during program execution.
- Only requirement of dynamic scheduling is that the worker **must** stay on its host core.

- 1 Background
- 2 Problem Statement
- 3 Related Work
- 4 Contributions
- 5 Motivation
- 6 Insights
- 7 Implementation**
 - Order of Immediacy
 - Tempo-Frequency Mapping
 - Worker-Core Mapping
 - **Tempo Setting at Idle**
 - Overheads
- 8 Experimental Methodology
- 9 Results
- 10 Conclusion
- 11 COTTON Integration

Implementation

Tempo Setting of Idle Workers/Core

- HERMES does not adjust CPU frequencies when a worker becomes idle (when pop and steal both fail).

Implementation

Tempo Setting of Idle Workers/Core

- HERMES does not adjust CPU frequencies when a worker becomes idle (when pop and steal both fail).
- YIELD: The core is reallocated to another worker.

Implementation

Tempo Setting of Idle Workers/Core

- HERMES does not adjust CPU frequencies when a worker becomes idle (when pop and steal both fail).
- YIELD: The core is reallocated to another worker.
- YIELD happens very rarely in practice.

Implementation

API

Algorithm 2.1 Worker

```
w : WORKER
procedure SCHEDULE(w)
  loop
    t ← POP(w)
    if t == null then
      v = SELECT()
      t ← STEAL(v)
      if t == null then
        YIELD(w)
      else
        WORK(w, t)
      end if
    else
      WORK(w, t)
    end if
  end loop
end procedure
```

Algorithm 3.1 Worker

```
1: w : WORKER
2: procedure SCHEDULE(w)
3:   loop
4:     t ← POP(w)
5:     if t == null then
6:       w0 = w.next
7:       for w0 != null do
8:         UP(w0)
9:         w0 ← w0.next
10:      end for
11:      w.prev.next ← w.next
12:      w.next.prev ← w.prev
13:      w.next ← null
14:      w.prev ← null
15:      v = SELECT()
16:      t ← STEAL(v)
17:      if t == null then
18:        YIELD(w)
19:      else
20:        DOWN(w, v)
21:        if v.next != null then
22:          w.next ← v.next
23:          v.prev ← w.prev
24:        end if
25:        v.next ← w
26:        w.prev ← v
27:        WORK(w, t)
28:      end if
29:    else
30:      WORK(w, t)
31:    end if
32:  end loop
33: end procedure
```

Algorithm 2.3 Pop

w : WORKER

procedure POP(**w**)

w.T ← −

if **w**.H > **w**.T **then**

w.T++

 LOCK(**w**)

w.T ← −

if **w**.H > **w**.T **then**

w.T++

 UNLOCK(**w**)

return null

end if

end if

 UNLOCK(**w**)

return **w**.DQ[**w**.T]

end procedure

Algorithm 3.4 Pop

w : WORKER

procedure POP(**w**)

w.T ← −

if **w**.H > **w**.T **then**

w.T++

 LOCK(**w**)

w.T ← −

if **w**.H > **w**.T **then**

w.T++

 UNLOCK(**w**)

return null

end if

end if

 UNLOCK(**w**)

if **w**.T - **w**.H < **w**.thld[**w**.S] **then**

if **w**.S > 0 **then**

if **w**.prev != **null** **then**

w.S ← −

 DOWN(**w**)

end if

end if

end if

return **w**.DQ[**w**.T]

end procedure

Implementation

API

Algorithm 2.4 Steal

```
v : WORKER // victim
procedure STEAL(v)
  LOCK(v)
  v.H++
  if v.H > v.T then
    v.H--
    UNLOCK(v)
    return null
  end if
  UNLOCK(v)
  return v.DQ[v.H]
end procedure
```

Algorithm 3.5 Steal

```
v : WORKER // victim
procedure STEAL(v)
  LOCK(v)
  v.H++
  if v.H > v.T then
    v.H--
    UNLOCK(v)
    return null
  end if
  UNLOCK(v)
  if v.T - v.H < v.thld[v.S] then
    if v.S > 0 then
      if v.prev != null then
        v.S--
        DOWN(v)
      end if
    end if
  end if
  return v.DQ[v.H]
end procedure
```

Implementation

API

Algorithm 2.2 Push

```
w : WORKER
t : TASK
procedure PUSH(w,t)
  w.T++
  w.DQ[w.T]  $\leftarrow$  t
end procedure
```

Algorithm 3.3 Push

```
w : WORKER
K: number of thresholds
t: TASK
procedure PUSH(w, t)
  w.T++
  w.DQ[w.T]  $\leftarrow$  t
  if w.T - w.H > w.thld[w.S] then
    if w.S < K-1 then
      w.S++
      UP(w)
    end if
  end if
end procedure
```

- 1 Background
- 2 Problem Statement
- 3 Related Work
- 4 Contributions
- 5 Motivation
- 6 Insights
- 7 Implementation**
 - Order of Immediacy
 - Tempo-Frequency Mapping
 - Worker-Core Mapping
 - Tempo Setting at Idle
 - Overheads
- 8 Experimental Methodology
- 9 Results
- 10 Conclusion
- 11 COTTON Integration

① DVFS switching cost

Implementation

Overheads

- 1 DVFS switching cost
- 2 Online profiling of workload threshold.

Implementation

Overheads

- 1 DVFS switching cost
- 2 Online profiling of workload threshold.
- 3 Affinity setting in dynamic scheduling.

Outline

1 Background

2 Problem Statement

3 Related Work

4 Contributions

5 Motivation

6 Insights

7 Implementation

8 Experimental Methodology

- Benchmarks
- System Specs
- Energy Consumption

9 Results

10 Conclusion

11 COTTON Integration

- 1 Background
- 2 Problem Statement
- 3 Related Work
- 4 Contributions
- 5 Motivation
- 6 Insights
- 7 Implementation
- 8 Experimental Methodology
 - Benchmarks
 - System Specs
 - Energy Consumption
- 9 Results
- 10 Conclusion
- 11 COTTON Integration

Experimental Methodology

Benchmarks

- K-Nearest Neighbors (`KNN`) uses pattern recognition methods to classify objects based on closest training examples in the feature space.
- Sparse Triangle Intersection (`Ray`) calculates for each ray the first triangle it intersects given a set of triangles contained in 3D bounding box and set of rays to penetrate.
- Integer Sort (`Sort`) is a parallel implementation of Radix Sort.
- Comparison Sort (`Compare`) is similar to `Sort` but uses sample sort.
- Convex Hull (`Hull`) is a computational geometry benchmark.

Experimental Methodology

Benchmarks

- K-Nearest Neighbors (`KNN`) uses pattern recognition methods to classify objects based on closest training examples in the feature space.
- Sparse Triangle Intersection (`Ray`) calculates for each ray the first triangle it intersects given a set of triangles contained in 3D bounding box and set of rays to penetrate.
- Integer Sort (`Sort`) is a parallel implementation of Radix Sort.
- Comparison Sort (`Compare`) is similar to `Sort` but uses sample sort.
- Convex Hull (`Hull`) is a computational geometry benchmark.
- We use `KNN` and `Ray` for our experiments.

- 1 Background
- 2 Problem Statement
- 3 Related Work
- 4 Contributions
- 5 Motivation
- 6 Insights
- 7 Implementation
- 8 Experimental Methodology**
 - Benchmarks
 - **System Specs**
 - Energy Consumption
- 9 Results
- 10 Conclusion
- 11 COTTON Integration

- Conducted experiments on two different systems to test effectiveness of the approach.
 - System A: 2 x 16-core AMD Opteron 6378 processors (Piledriver microarchitecture) with 64GB DDR3 1600 memory
 - System B: 8-core AMD FX-8150 processor (Bulldozer microarchitecture) with 16GB DDR3 1600 memory
- Both systems run Debian 3.2.46-1 x86-64 Linux (kernel 3.2.0-4-amd64) and support 5 different CPU frequencies.

- Conducted experiments on two different systems to test effectiveness of the approach.
 - System A: 2 x 16-core AMD Opteron 6378 processors (Piledriver microarchitecture) with 64GB DDR3 1600 memory
 - System B: 8-core AMD FX-8150 processor (Bulldozer microarchitecture) with 16GB DDR3 1600 memory
- Both systems run Debian 3.2.46-1 x86-64 Linux (kernel 3.2.0-4-amd64) and support 5 different CPU frequencies.
- Piledriver / Bulldozer were chosen as they supported *multiple clock domains*.

Outline

1 Background

2 Problem Statement

3 Related Work

4 Contributions

5 Motivation

6 Insights

7 Implementation

8 Experimental Methodology

- Benchmarks
- System Specs
- Energy Consumption

9 Results

10 Conclusion

11 COTTON Integration

Experimental Methodology

Energy Consumption

- Measured through current meters over power supply lines to the CPU module.
- Data converted to NI DAQ and collected by NI LabVIEW SignalExpress with 100 samples per second.
- Supply voltage stable at 12V, so energy consumption computed as sum of current samples multiplied by 12×0.01

Outline

1 Background

2 Problem Statement

3 Related Work

4 Contributions

5 Motivation

6 Insights

7 Implementation

8 Experimental Methodology

9 Results

- Overall Results
- Relative Effectiveness
- Effect of Frequency Selection
- N-Frequency Tempo Control
- Static vs Dynamic Scheduling
- Naive Frequency Scaling

10 Conclusion

11 COTTON Integration

Outline

1 Background

2 Problem Statement

3 Related Work

4 Contributions

5 Motivation

6 Insights

7 Implementation

8 Experimental Methodology

9 Results

- Overall Results
 - Relative Effectiveness
 - Effect of Frequency Selection
 - N-Frequency Tempo Control
 - Static vs Dynamic Scheduling
 - Naive Frequency Scaling

10 Conclusion

11 COTTON Integration

Results

Overall Results

- On system A, experiments were conducted using 2, 4, 8, 16 workers.
- On system B, experiments were conducted using 2, 3, 4 workers.



Results

Overall Results

- On system A, experiments were conducted using 2, 4, 8, 16 workers.
- On system B, experiments were conducted using 2, 3, 4 workers.
- In both systems, HERMES averages 11-12% energy savings over 3-4% performance loss.

Results

Overall Results



Results

Overall Results

- Energy Delay Product(EDP): Product of energy consumption and execution time.

Results

Overall Results

- Energy Delay Product(EDP): Product of energy consumption and execution time.
- EDP is often used as an indicator for demonstrating the energy/performance tradeoff.



Results

Overall Results

- Energy Delay Product(EDP): Product of energy consumption and execution time.
- EDP is often used as an indicator for demonstrating the energy/performance tradeoff.
- In both System A and System B, the average normalized EDP is about 0.92.

Results

Overall Results



Outline

1 Background

2 Problem Statement

3 Related Work

4 Contributions

5 Motivation

6 Insights

7 Implementation

8 Experimental Methodology

9 Results

- Overall Results
- **Relative Effectiveness**
- Effect of Frequency Selection
- N-Frequency Tempo Control
- Static vs Dynamic Scheduling
- Naive Frequency Scaling

10 Conclusion

11 COTTON Integration

Results

Relative Effectiveness

- We also run benchmarks with only one of the two tempo-control strategies enabled.

Results

Relative Effectiveness

- We also run benchmarks with only one of the two tempo-control strategies enabled.
- Data are normalized w.r.t to the unified HERMES algorithm.



Results

Relative Effectiveness

- We also run benchmarks with only one of the two tempo-control strategies enabled.
- Data are normalized w.r.t to the unified HERMES algorithm.
- The following set of figures show the complementary nature of the two strategies.



Results

Relative Effectiveness



Results

Relative Effectiveness



1 Background

2 Problem Statement

3 Related Work

4 Contributions

5 Motivation

6 Insights

7 Implementation

8 Experimental Methodology

9 Results

- Overall Results
- Relative Effectiveness
- **Effect of Frequency Selection**
- N-Frequency Tempo Control
- Static vs Dynamic Scheduling
- Naive Frequency Scaling

10 Conclusion

11 COTTON Integration

Results

Effect of Frequency Selection

- We experimentally evaluate the effects of different frequency mapping strategies.

Results

Effect of Frequency Selection

- We experimentally evaluate the effects of different frequency mapping strategies.
- We only consider 2-frequency tempo control.
 - Fastest tempo is mapped to the first frequency.
 - All other tempos are mapped to the second frequency.



Results

Effect of Frequency Selection

- We experimentally evaluate the effects of different frequency mapping strategies.
- We only consider 2-frequency tempo control.
 - Fastest tempo is mapped to the first frequency.
 - All other tempos are mapped to the second frequency.
- We fix the frequency for the fast tempo 2.4Ghz for System A and 3.6GHz for System B.
 - Experiment with different settings for the slow tempo.

Results

Effect of Frequency Selection



Results

Effect of Frequency Selection

- Two scenarios:
 - 1 Higher frequency for slow tempo.
 - 2 Lower frequency for slow tempo.

Results

Effect of Frequency Selection

- Two scenarios:
 - 1 Higher frequency for slow tempo.
 - 2 Lower frequency for slow tempo.
- Optimal combination:
 - Frequency of slow tempo should be 60% of that of fast tempo.

1 Background

2 Problem Statement

3 Related Work

4 Contributions

5 Motivation

6 Insights

7 Implementation

8 Experimental Methodology

9 Results

- Overall Results
- Relative Effectiveness
- Effect of Frequency Selection
- **N-Frequency Tempo Control**
- Static vs Dynamic Scheduling
- Naive Frequency Scaling

10 Conclusion

11 COTTON Integration



Results

N-Frequency Tempo Control

- We study how the number of frequencies impact the results.

Results

N-Frequency Tempo Control



Results

N-Frequency Tempo Control

- A 3-frequency tempo control can sometimes incur less loss on performance.

Results

N-Frequency Tempo Control

- A 3-frequency tempo control can sometimes incur less loss on performance.
- But the 2-frequency tempo control has a slight edge on energy savings.

Results

N-Frequency Tempo Control

- A 3-frequency tempo control can sometimes incur less loss on performance.
- But the 2-frequency tempo control has a slight edge on energy savings.
- Maybe due to lesser overhead on DVFS.

1 Background

2 Problem Statement

3 Related Work

4 Contributions

5 Motivation

6 Insights

7 Implementation

8 Experimental Methodology

9 Results

- Overall Results
- Relative Effectiveness
- Effect of Frequency Selection
- N-Frequency Tempo Control
- **Static vs Dynamic Scheduling**
- Naive Frequency Scaling

10 Conclusion

11 COTTON Integration


Results

Static vs Dynamic Scheduling

- We study the effectiveness of HERMES under static scheduling and dynamic scheduling.

Results

N-Frequency Tempo Control



Results

Static vs Dynamic Scheduling

- Following figures show a more detailed analysis.
- Time series plot of power consumption.



Results

Static vs Dynamic Scheduling

- Following figures show a more detailed analysis.
- Time series plot of power consumption.
- The shape of the time series are dependent on the type benchmarks and their settings (such as number of workers).



Results

N-Frequency Tempo Control



Results

N-Frequency Tempo Control



Results

Static vs Dynamic Scheduling

- For each benchmark with the same number of workers, the executions of static scheduling vs. dynamic scheduling display similar patterns.

Results

Static vs Dynamic Scheduling

- For each benchmark with the same number of workers, the executions of static scheduling vs. dynamic scheduling display similar patterns.
- Dynamic scheduling incurs a slightly higher level of energy consumption.

Results

Static vs Dynamic Scheduling

- For each benchmark with the same number of workers, the executions of static scheduling vs. dynamic scheduling display similar patterns.
- Dynamic scheduling incurs a slightly higher level of energy consumption.
- Maybe due to the overhead of setting/re-setting the affinity of workers.

Outline

1 Background

2 Problem Statement

3 Related Work

4 Contributions

5 Motivation

6 Insights

7 Implementation

8 Experimental Methodology

9 Results

- Overall Results
- Relative Effectiveness
- Effect of Frequency Selection
- N-Frequency Tempo Control
- Static vs Dynamic Scheduling
- Naive Frequency Scaling

10 Conclusion

11 COTTON Integration

Results

Naive Frequency Scaling

- Finally, we investigate the impact of naively applying frequency scaling.



Results

Naive Frequency Scaling

- Finally, we investigate the impact of naively applying frequency scaling.
- CPU frequencies are naively scaled down at a fixed frequency throughout the program execution.

Results

Naive Frequency Scaling



Outline

- 1 Background
- 2 Problem Statement
- 3 Related Work
- 4 Contributions
- 5 Motivation
- 6 Insights
- 7 Implementation
- 8 Experimental Methodology
- 9 Results
- 10 Conclusion**
- 11 COTTON Integration

Conclusion

- Proposed HERMES – a novel and practical solution for improving energy efficiency of work-stealing applications.

Conclusion

- Proposed HERMES – a novel and practical solution for improving energy efficiency of work-stealing applications.
- Solution addresses the problem through judicious tempo control over workers, guided by a unified algorithm.

Conclusion

- Proposed HERMES – a novel and practical solution for improving energy efficiency of work-stealing applications.
- Solution addresses the problem through judicious tempo control over workers, guided by a unified algorithm.
- Solution is scalable as it only makes minor changes to existing work-stealing runtimes.

Conclusion

- Proposed HERMES – a novel and practical solution for improving energy efficiency of work-stealing applications.
- Solution addresses the problem through judicious tempo control over workers, guided by a unified algorithm.
- Solution is scalable as it only makes minor changes to existing work-stealing runtimes.
- HERMES is a language-level solution without the need of characterizing program executions on a per-application basis.

Outline

- 1 Background
- 2 Problem Statement
- 3 Related Work
- 4 Contributions
- 5 Motivation
- 6 Insights
- 7 Implementation
- 8 Experimental Methodology
- 9 Results
- 10 Conclusion
- 11 COTTON Integration

COTTON Integration

Work-First vs Help-First

- The COTTON runtime is help-first work-stealing runtime.
 - HERMES follows the work-first principle.

COTTON Integration

Work-First vs Help-First

- The COTTON runtime is help-first work-stealing runtime.
 - HERMES follows the work-first principle.
- The sequence of tasks pushed in the deque remains the same.

COTTON Integration

Work-First vs Help-First

- The COTTON runtime is help-first work-stealing runtime.
 - HERMES follows the work-first principle.
- The sequence of tasks pushed in the deque remains the same.
- No scenario yet identified where HERMES may show rogue behaviour in a help-first environment.

COTTON Integration

New Worker structure

- Member variables
 - Deque pointer
 - Linked list to maintain victim-thief and immediacy relationships
 - Frequency thresholds
- Member functions
 - Up(worker)
 - Down(worker)
 - Down(worker, victim)

COTTON Integration

Modified methods

- Worker routine
- Deque Push
- Deque Pop
- Deque Steal
- Deque Push

COTTON Integration

Modified methods

- Worker routine
- Deque Push
- Deque Pop
- Deque Steal
- Deque Push
- The corresponding modifications have been covered in the Implementation section.